

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: data=pd.read_csv('BMW_Car_Sales_Classification.csv')
```

```
In [3]: #data=data.drop('Sales_Classification',axis=1)
```

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Model                 50000 non-null object  
 1   Year                 50000 non-null int64   
 2   Region               50000 non-null object  
 3   Color                50000 non-null object  
 4   Fuel_Type            50000 non-null object  
 5   Transmission         50000 non-null object  
 6   Engine_Size_L        50000 non-null float64  
 7   Mileage_KM           50000 non-null int64   
 8   Price_USD            50000 non-null int64   
 9   Sales_Volume         50000 non-null int64   
10   Sales_Classification 50000 non-null object  
dtypes: float64(1), int64(4), object(6)
memory usage: 4.2+ MB
```

```
In [5]: data= pd.get_dummies(data,columns=['Sales_Classification','Transmission','Fuel_Type'])
```

```
In [6]: data.head()
```

```
Out[6]:
```

	Year	Engine_Size_L	Mileage_KM	Price_USD	Sales_Volume	Sales_Classification_High	Sales_Classification_Low
--	------	---------------	------------	-----------	--------------	---------------------------	--------------------------

0	2016	3.5	151748	98740	8300	True	False
1	2013	1.6	121671	79219	3428	False	False
2	2022	4.5	10991	113265	6994	False	False
3	2024	1.7	27255	60971	4047	False	False
4	2020	2.1	122131	49898	3080	False	False


5 rows × 36 columns



```
In [7]: data.describe()
```


	Year	Engine_Size_L	Mileage_KM	Price_USD	Sales_Volume
<b>count</b>	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000
<b>mean</b>	2017.015700	3.247180	100307.203140	75034.600900	5067.514680
<b>std</b>	4.324459	1.009078	57941.509344	25998.248882	2856.767125
<b>min</b>	2010.000000	1.500000	3.000000	30000.000000	100.000000
<b>25%</b>	2013.000000	2.400000	50178.000000	52434.750000	2588.000000
<b>50%</b>	2017.000000	3.200000	100388.500000	75011.500000	5087.000000
<b>75%</b>	2021.000000	4.100000	150630.250000	97628.250000	7537.250000
<b>max</b>	2024.000000	5.000000	199996.000000	119998.000000	9999.000000

```
In [8]: plt.hist(data['Price_USD'])
```



```
In [9]: sns.heatmap(data.corr(),annot=True)
```

```
Out[9]: <Axes: >
```



```
In [10]: #sns.pairplot(data)
```

```
In [11]: x=data.drop('Price_USD',axis=1)
y=data['Price_USD']
```

```
In [12]: x
```

Out[12]:

	Year	Engine_Size_L	Mileage_KM	Sales_Volume	Sales_Classification_High	Sales_Cla
0	2016	3.5	151748	8300		True
1	2013	1.6	121671	3428		False
2	2022	4.5	10991	6994		False
3	2024	1.7	27255	4047		False
4	2020	2.1	122131	3080		False
...	...	...	...	...		...
49995	2014	4.6	151030	8182		True
49996	2023	4.2	147396	9816		True
49997	2010	4.5	174939	8280		True
49998	2020	3.8	3379	9486		True
49999	2020	3.3	171003	1764		False

50000 rows × 35 columns

In [13]:

y

Out[13]:

0	98740
1	79219
2	113265
3	60971
4	49898
...	
49995	42932
49996	48714
49997	46126
49998	58566
49999	77492

Name: Price\_USD, Length: 50000, dtype: int64

In [14]:

from sklearn.model\_selection import train\_test\_split

In [15]:

x\_train,x\_test,y\_train,y\_test=train\_test\_split(x,y,test\_size=0.3,random\_state=23)

In [16]:

from sklearn.linear\_model import LinearRegression

In [17]:

lm=LinearRegression()

In [18]:

lm.fit(x\_train,y\_train)

Out[18]:

▼ LinearRegression ⓘ ?  
► Parameters

In [19]: `lm.intercept_`

Out[19]: `np.float64(-11224.909669668006)`

In [20]: `lm.coef_`

Out[20]: `array([ 4.24126815e+01, 1.89967550e+02, -2.84274649e-03, 3.18948016e-02,  
 -1.05206128e+02, 1.05206128e+02, 1.16595260e+02, -1.16595260e+02,  
 -9.06046530e+01, 3.95268207e+02, -3.05013646e+02, 3.50092234e-01,  
 -1.98122556e+02, 4.34024316e+02, -1.33125912e+01, -1.93396990e+02,  
 -4.91516709e+02, 4.62324530e+02, -1.33310722e+02, 6.52763473e+02,  
 -1.61882107e+02, -2.85927177e+02, -1.54395104e+02, 8.27516374e+01,  
 2.94138016e+02, -1.07911494e+01, 3.95824290e+02, -1.79699298e+02,  
 -6.90159342e+02, 5.88820823e+02, -9.13124947e+01, -2.22120044e+02,  
 -4.53282952e+02, 1.42682954e+02, 2.25899197e+02])`

In [21]: `c=pd.DataFrame(lm.coef_,x.columns,columns=['Price_USD'])`

In [22]: `c`

Out[22]:

	Price_USD
Year	42.412682
Engine_Size_L	189.967550
Mileage_KM	-0.002843
Sales_Volume	0.031895
Sales_Classification_High	-105.206128
Sales_Classification_Low	105.206128
Transmission_Automatic	116.595260
Transmission_Manual	-116.595260
Fuel_Type_Diesel	-90.604653
Fuel_Type_Electric	395.268207
Fuel_Type_Hybrid	-305.013646
Fuel_Type_Petrol	0.350092
Color_Black	-198.122556
Color_Blue	434.024316
Color_Grey	-13.312591
Color_Red	-193.396990
Color_Silver	-491.516709
Color_White	462.324530
Region_Africa	-133.310722
Region_Asia	652.763473
Region_Europe	-161.882107
Region_Middle East	-285.927177
Region_North America	-154.395104
Region_South America	82.751637
Model_3 Series	294.138016
Model_5 Series	-10.791149
Model_7 Series	395.824290
Model_M3	-179.699298
Model_M5	-690.159342
Model_X1	588.820823

	Price_USD
<b>Model_X3</b>	-91.312495
<b>Model_X5</b>	-222.120044
<b>Model_X6</b>	-453.282952
<b>Model_i3</b>	142.682954
<b>Model_i8</b>	225.899197

In [23]: `pr=lm.predict(x_test)`

In [24]: `pr`


Out[24]: `array([73500.79639859, 74633.70418724, 74779.94638861, ...,  
73620.55249048, 75150.90617265, 74552.66215489], shape=(15000,))`

In [25]: `y_test`

Out[25]: `49466 104395  
11621 108196  
39058 86604  
10033 101480  
22076 59951  
...  
35962 71658  
1710 41403  
37523 119878  
14167 92089  
21600 30170  
Name: Price_USD, Length: 15000, dtype: int64`

In [26]: `plt.scatter(x=pr,y=y_test)`

Out[26]: `<matplotlib.collections.PathCollection at 0x2d72183c500>`



```
In [27]: from sklearn import metrics
```

```
In [28]: metrics.mean_absolute_error(y_test,pr)
```

```
Out[28]: 22660.005257327874
```

```
In [29]: metrics.mean_squared_error(y_test,pr)
```

```
Out[29]: 681578803.0890763
```

```
In [30]: np.sqrt(metrics.mean_squared_error(y_test,pr))
```

```
Out[30]: np.float64(26107.064237272567)
```

```
In [31]: metrics.r2_score(y_test,pr)
```

```
Out[31]: -0.0012332612175953717
```

```
In [32]: sns.displot(y_test-pr,bins=25)
```

```
Out[32]: <seaborn.axisgrid.FacetGrid at 0x2d721479850>
```



